

Building Smarter Agents with LangGraph Tools, Memory & Workflows

A quick guide to using tools and memory in your agent workflows



RAJNEESH JHA

JUN 05, 2025 • PAID



In my recent articles, I've been diving into various tools and libraries for developing multi-agent systems.

We started with the **OpenAI Agent SDK**, where I shared how to build your own Research Agent:

[🔗 Build Your Own Deep Research Agent using OpenAI's Agent SDK](#)

Then we explored **CrewAI**, a powerful framework for creating collaborative AI agents.

[🔗 Build Your First AI Coding Buddy with CrewAI](#)

Today, let's take a look at another exciting and widely adopted framework — **LangGraph**. This graph-based library is designed for building complex, structured agent workflows with ease. In the next few minutes, we'll see how to use tools and memory with LangGraph.




Image Source: <https://realpython.com/langgraph-python/>

1. What is LangGraph ?

LangGraph is a stateful, graph-based framework designed to build complex, multi-agent applications using LLMs. Unlike traditional linear workflows, LangGraph leverages directed graphs where each node represents an agent or task, and edges define their interactions. This structure facilitates dynamic, iterative, and memory-aware agentic systems.

Core Features of LangGraph:

- **Graph-Oriented Workflow:** Design flexible, branching workflows with loops and conditionals, enabling complex agent interactions.
- **Stateful Execution:** Automatically persist agent states at each step, allowing error recovery, time travel, and human-in-the-loop interventions.
- **Integrated Memory:** Support for both short-term (session-based) and long-term (persistent across sessions) memory, enabling agents to recall and adapt to user interactions.

- **Human-in-the-Loop Control:** Pause execution to seek human approval or modification before proceeding, enhancing oversight and control .
- **Seamless Integration with LangChain:** While built on LangChain, LangGraph can function independently, offering flexibility in deployment.
- **Open-Source and Extensible:** Released under the MIT license, LangGraph is free to use and can be extended to meet specific needs.

2. Key Terminology in LangGraph:

1. The **State** serves as the central data structure in LangGraph, encapsulating all relevant information during the execution of the workflow. It can be as simple as a dictionary or a more structured model using tools like Pydantic. Each node in the graph receives and updates this state, ensuring continuity and context throughout the process.

When defining the state in LangGraph, two important concepts from LangChain play a crucial role:

- **Annotated from Python's typing module** is used to attach metadata to state fields. This metadata usually specifies a **reducer function**, which dictates how updates to a particular state field should be merged with existing data.
- **Reducers** are functions responsible for determining how updates to individual state fields are applied. Each field can have its own reducer, allowing for customized merging behaviour. If no reducer is specified, the default behaviour is to overwrite the existing value with the new update.

```
from typing import Annotated
from typing_extensions import TypedDict
from operator import add
```

```
class State(TypedDict):
    messages: Annotated[list[str], add]
```

In this example, the `messages` field is a list of strings, and the `add` function from the `operator` module is used as a reducer. This means that when new messages are added, they will be concatenated to the existing list rather than replacing it.

2. Nodes are the fundamental units of computation within LangGraph. Each node represents a specific task or action, such as:

- Processing user input
- Interacting with external APIs
- Making decisions based on conditions
- Generating responses using LLMs

Nodes are typically implemented as Python functions that accept the current state as input and return an updated state.

3. Edges define the flow of execution between nodes, determining the sequence in which tasks are performed. There are different types of edges:

- **Normal Edges:** Establish a straightforward, unidirectional connection between nodes.
- **Conditional Edges:** Introduce decision-making into the workflow, directing flow based on specific conditions or logic.

These edges enable the creation of complex workflows with branching logic and

4. Unlike traditional DAGs, LangGraph supports **cycles**, allowing nodes to revisit previous steps. This feature is particularly useful for tasks that require iterative refinement or repeated evaluation, such as debates between agents or recursive problem-solving.

5. LangGraph provides built-in persistence mechanisms to save the state of the workflow at various points. This capability enables:

- **Checkpointing:** Saving snapshots of the graph's state to resume execution later.

- **Human-in-the-Loop Interactions:** Allowing manual intervention during the workflow.
- **Fault Tolerance:** Recovering from errors by restoring to a previous state.

Persistence is achieved through check pointers and storage systems that manage across different threads and sessions.

6. In LangGraph, a **thread** represents an individual session or conversation. Each thread maintains its own state, enabling personalized interactions and context retention across multiple sessions. This is particularly beneficial for applications chatbots or virtual assistants that require memory of past interactions.

7. LangGraph includes a **storage** system that allows for the saving and retrieval c across different threads. This feature supports:

- **Cross-Thread Persistence:** Maintaining information across multiple session
- **Flexible Name spacing:** Organizing data for different users or contexts.
- **Document Storage:** Saving data in JSON format for easy manipulation and retrieval.

The storage system facilitates the creation of long-term memories and knowledg bases.

In the next few sections, let's create some tools and integrate with agents.

3. Tool Creation:

Serper Web-Search Tool: Serper is a web search API tool that provides real-time Google-like search results. It's often used in AI and agent frameworks like LangG to allow language models to access up-to-date information from the web.

Explore more here: <https://serper.dev/>

```

from langchain_community.utilities import GoogleSerperAPIWrapper

serper = GoogleSerperAPIWrapper()
serper.run("What is the capital of India?")

from langchain.agents import Tool

tool_search = Tool(
    name="search",
    func=serper.run,
    description=(
        "A web search tool powered by Serper.dev. Use this tool when\n"
        "need to retrieve real-time "\n"
        "information from the internet, such as current events, rece\n"
        "news, public data, or facts that "\n"
        "are not part of the model's training knowledge. Ideal for\n"
        "answering up-to-date queries, verifying information, "\n"
        "or supplementing static knowledge with fresh content from l\n"
        "web sources."
    )
)

tool_search.invoke("What is the capital of India?")

```

Push Over Tool: Pushover is a simple and reliable push notification service that allows you to send real-time alerts to your phone, tablet, or desktop. It's widely used by developers and system administrators to notify users or themselves of important events like server errors, completed tasks, or AI agent actions.

Explore more here: <https://pushover.net/>

```

import os
import requests
from langchain.agents import Tool

# Load credentials from environment variables
pushover_token = os.getenv("PUSHOVER_TOKEN")
pushover_user = os.getenv("PUSHOVER_USER")
pushover_url = "https://api.pushover.net/1/messages.json"

```

```

def push(text: str):
    """Send a push notification using the Pushover API."""
    if not pushover_token or not pushover_user:
        raise ValueError("Pushover credentials are missing. Check
environment variables.")

    payload = {
        "token": pushover_token,
        "user": pushover_user,
        "message": text
    }

    response = requests.post(pushover_url, data=payload)

    if response.status_code != 200:
        raise Exception(f"Failed to send push notification:
{response.text}")

tool_push = Tool(
    name="send_push_notification",
    func=push,
    description=(
        "Useful for sending real-time push notifications to the user
the Pushover API. "
        "Use this when you need to alert the user with important
messages or status updates."
    )
)

tool_push.invoke("Hello, Good People !")
# Check the notification in your pushover app on your phone.

```

Now that the tools are set up, let's build a basic graph with nodes powered by an LLM that can use these tools.

4. Define Workflow:

In this section, we'll create a basic **LangGraph** powered by an **LLM** that can integrate with external tools. We've already defined two tools:

- A **web search tool** (using Serper) for real-time internet queries
- A **push notification tool** (using Pushover) to send alerts

Now, we'll build a simple graph where:

- An LLM node processes user input
- If tool use is needed, the graph routes to a tool node
- After the tool is invoked, control returns to the LLM to decide what to do ne

```
from typing import Annotated
from langgraph.graph import StateGraph, START, END
from langgraph.graph.message import add_messages
from dotenv import load_dotenv
from IPython.display import Image, display
import gradio as gr
from langgraph.prebuilt import ToolNode, tools_condition
import requests
import os
from langchain_openai import ChatOpenAI
from typing import TypedDict

load_dotenv(override=True)

# Create a list of Tools
tools = [tool_search, tool_push]

# Define State
class State(TypedDict):
    messages: Annotated[list, add_messages]

# Start the Graph Building process
graph_builder = StateGraph(State)

# Define the LLM to be used
llm = ChatOpenAI(model="gpt-4o")
llm_with_tools = llm.bind_tools(tools)

# Create a Node
def chatbot(state: State):
```



```

        return {"messages": [llm_with_tools.invoke(state["messages"])]}

# Add Nodes
graph_builder.add_node("chatbot", chatbot)
graph_builder.add_node("tools", ToolNode(tools=tools))

# If any tool is called, we return to the chatbot to decide the next
step
graph_builder.add_conditional_edges( "chatbot", tools_condition,
"tools")
graph_builder.add_edge("tools", "chatbot")

graph_builder.add_edge(START, "chatbot")

# Compile the Graph
graph = graph_builder.compile()
display(Image(graph.get_graph().draw_mermaid_png()))

```

We can use a simple [Gradio](#) Chat interface to test this graph quickly:

Note: Before running, make sure you have below things in environment variables ([LangSmith](#) is optional, if you want to trace the agent in a better way.)

- OPENAI_API_KEY
- LANGSMITH_TRACING=true
- LANGSMITH_ENDPOINT
- LANGSMITH_API_KEY
- LANGSMITH_PROJECT
- SERPER_API_KEY
- PUSHOVER_USER
- PUSHOVER_TOKEN
- PUSHOVER_URL

```

def chat(user_input: str, history):
    result = graph.invoke({"messages": [{"role": "user", "content":

```

```

user_input}}})
    return result["messages"][-1].content

gr.ChatInterface(chat, type="messages").launch()

```

5. Adding Memory:

You might be wondering, *“If LangGraph maintains and updates state between nodes, do we still need memory?”*

LangGraph processes workflows in **units called “super-steps”**.

- A **super-step** is a single pass through the graph running one or more nodes (possibly in parallel), based on the current state.
- Each time you call `graph.invoke()`, you're running **one super-step**.
- **Reducers** handle state updates within this super-step they control how data merged or modified while the graph is running.

Once the super-step completes, the graph’s in-memory state is gone unless you explicitly save it.

This means:

- LangGraph manages **short-term state** during execution (via the state object reducers)
- But it doesn’t **automatically persist memory** between interactions or sessions

That’s where **checkpointing and memory storage** come in.

To retain state between graph invocations, LangGraph offers **persistence mechanisms** like:

- **Checkpointing:** Save the state after each super-step so it can be resumed or referenced later

- **Storage Backends:** For managing long-term memory or user-specific conversations across sessions

This enables:

- Stateful, ongoing conversations
- Time travel (restoring to previous states)
- Human-in-the-loop interventions
- Memory-aware agents

Let's now enhance the LangGraph with **memory support** using a **checkpointing system**. This allows the agent to retain and recall conversation history across sup steps (individual runs of the graph).

```
from langgraph.checkpoint.memory import MemorySaver

# MemorySaver is a built-in checkpointing utility.
# It stores the state of the graph after each run, allowing recovery
# continuation later.

memory = MemorySaver()

# This means the graph automatically saves state after each interact
# (super
# step).
graph = graph_builder.compile(checkpointer=memory)
display(Image(graph.get_graph().draw_mermaid_png()))

# This is the session ID or conversation thread ID. Think of it like
# chatroom. All messages in the same thread share memory.
config = {"configurable": {"thread_id": "1"}}

def chat(user_input: str, history):
    result = graph.invoke({"messages": [{"role": "user", "content":
user_input}]}, config=config)
    return result["messages"][-1].content
```

```

gr.ChatInterface(chat, type="messages").launch()

graph.get_state(config)

list(graph.get_state_history(config))

# By specifying a checkpoint_id, you can restore the graph to a prev
point # in time.
# config = {"configurable": {"thread_id": "1", "checkpoint_id": ...}}
# graph.invoke(None, config=config)

```

Persisting Memory:

Previously, we used **MemorySaver** for temporary in-memory persistence (RAM). We can also use **SQLite**, which gives LangGraph application **disk-based, long-term memory**. This means the agent can remember past interactions **even after restarting the program**.

```

import sqlite3
from langgraph.checkpoint.sqlite import SqliteSaver

# SqliteSaver is LangGraph's utility to store state checkpoints in a
table.
# This allows for long-term persistence, cross-session recall, and s
# restoration from disk.

db_path = "memory.db"
conn = sqlite3.connect(db_path, check_same_thread=False)
sql_memory = SqliteSaver(conn)

graph = graph_builder.compile(checkpointer=sql_memory)
display(Image(graph.get_graph().draw_mermaid_png()))

# Thread ID "3" defines a unique session. Each session has its own
memory.
# This allows the agent to carry on context from earlier messages ac
# restarts.

```

```
config = {"configurable": {"thread_id": "3"}}

def chat(user_input: str, history):
    result = graph.invoke({"messages": [{"role": "user", "content":
user_input}]}, config=config)
    return result["messages"][-1].content

gr.ChatInterface(chat, type="messages").launch()
```

This should give the agent a persistent memory that remains intact even after a kernel restart. Try mentioning your name, then restart the kernel the agent should still remember it

6. Conclusion:

In this post, we went through the fundamentals of LangGraph a powerful, graph framework for building structured, multi-agent AI workflows.

We explored:

- What LangGraph is and how it differs from linear agent frameworks
- Key features like stateful execution, memory integration, and human-in-the-loop control
- Core concepts such as State, Nodes, Edges, and how they work together to create dynamic workflows
- Building a simple agent that is integrated with multiple tools
- Adding memory to that agent and persisting the memory

As agentic systems continue to evolve, frameworks like LangGraph will play a key role in making them more reliable, transparent, and customizable. Stay tuned! In the next few posts, we'll dive into hands-on examples of building real-world agent workflows with tools and memory using LangGraph.

Thanks for Reading. Have any questions or concerns? Feel free to connect, and stay tuned for more insightful content. Let's connect on LinkedIn:

<https://www.linkedin.com/in/rajneesh407/>

To Data & Beyond is a reader-supported publication. To receive new posts and support my work, consider becoming a free or paid subscriber.



8 Likes · 2 Restacks

← Previous

Next →



A guest post by

Rajneesh Jha

Interested in solving business problems using Machine Learning and AI. <https://www.linkedin.com/in/rajneesh407/>

Subscribe to Rajneesh Jha

Discussion about this post

Comments

Restacks



Write a comment...

© 2025 Youssef Hosni · [Privacy](#) · [Terms](#) · [Collection notice](#)
[Substack](#) is the home for great culture